

Financial Tweet Sentiment Classification

Multiclass Sentiment Analysis of Financial Social Media Text

Text Mining — MSc in Data Science and Advanced Analytics

NOVA Information Management School (NOVA IMS)

Group 33 · Academic Year 2025/2026 · June 2026

Abstract

This report presents a multiclass sentiment classification system for financial tweets, distinguishing three market stances: Bearish (0), Bullish (1), and Neutral (2). The dataset comprises 9,543 labelled training tweets with pronounced class imbalance (Neutral 64.7%, Bullish 20.1%, Bearish 15.1%; imbalance ratio 4.28:1) and substantial text corruption: 7.0% of tweets carry UTF-8 mojibake and 16.4% end in truncation artefacts. We systematically evaluate seven text preprocessing techniques, eleven feature representations — spanning sparse TF-IDF, Word2Vec, GloVe-Twitter, FinBERT, SBERT and Twitter-RoBERTa embeddings — and twelve model families under a common stratified cross-validation protocol (10-fold for all transformer fine-tuning, seed 42). Frozen-embedding classifiers tuned with Optuna plateau around macro-F1 0.80. A **transformer backbone search** over nine fine-tuned encoders shows that domain-matched pre-training dominates model size: a 110M-parameter **FinBERT** variant pre-adapted to financial Twitter text reaches an out-of-fold macro-F1 of **0.9082**, beating DeBERTa-v3-large (435M, 0.8889). As extra work, a weighted soft-vote **ensemble** of eight encoders — with weights found by exhaustive out-of-fold grid search — attains **0.9201**, a statistically significant gain over the best individual model (bootstrap 95% CI of the difference: [+0.007, +0.017]). To comply with the single-model delivery rule, we **distil** this ensemble into one FinBERT student (Hinton et al., 2015), our **submitted single-model solution** at OOF macro-F1 **0.9134**. We additionally show that a fine-tuned **GPT-2 decoder** classifies at macro-F1 0.7724 (whereas few-shot prompting collapses to the majority class), and that an offline **agentic workflow** adaptively routes tweets between VADER, LightGBM/SBERT and FinBERT with conversational explanations.

Keywords: Financial sentiment analysis, tweet classification, FinBERT fine-tuning, model selection, ensemble distillation, decoder classification, LangChain agent.

1 Introduction

Social media platforms — particularly Twitter/X — have become a primary channel through which retail and institutional investors express and discover market sentiment. Automated classification of financial tweets into Bearish, Bullish and Neutral stances enables real-time signal generation for algorithmic trading, risk monitoring and investor-sentiment tracking (Bollen et al., 2011; Antweiler & Frank, 2004).

This work addresses multiclass financial sentiment classification — a text-categorization problem in the classical sense of Sebastiani (2002) — under significant class imbalance (ratio 4.28:1). We explore the full NLP pipeline — from raw text through feature engineering to model selection and hyperparameter tuning — following the Text Mining curriculum at NOVA IMS. All models are trained

only on the provided training corpus. The remainder of the report is organised as follows: Section 2 explores the dataset; Section 3 analyses preprocessing; Section 4 describes feature engineering; Section 5 presents the classification models, including the backbone search and the fine-tuned encoder/decoder; Section 6 reports results and error analysis; Section 7 details the agentic workflow (Extra Challenge 2); Section 8 concludes.

Course alignment. Every stage follows the Text Mining curriculum: the NLP pipeline and text preprocessing (Lecture 1, Lab 1); n-grams, TF-IDF and distance metrics (Lecture 2-Extra); word2vec word embeddings and their semantic geometry (Lecture 2, Lab 2); attention and Transformer *encoders* (Lecture 4, Lab 4); and *decoder* LLMs, encoder-based classification and *Agentic AI* (Lecture 6, Lab 5; Alammari & Grootendorst, 2024).

2 Data Exploration

The dataset consists of 9,543 labelled training tweets and 2,388 unlabelled test tweets from financial Twitter discourse. Each tweet references one or more publicly-listed securities and expresses a market stance.

2.1 Class Distribution

Figure 1 illustrates the severe class imbalance. Neutral tweets dominate (64.7%), consistent with most financial discourse being non-directional. Bearish tweets (15.1%) are less frequent than Bullish (20.1%), reflecting a structural optimism bias in financial social media (Antweiler & Frank, 2004). To address imbalance, all models use `class_weight='balanced'` (or an inverse-frequency-weighted loss) and evaluation relies on **macro-F1** as the primary metric, which weights all three classes equally regardless of support.




Figure 1: Class distribution of the training set: absolute counts (left) and proportions (right).

2.2 Linguistic Characteristics

Tweets have a mean length of 12.2 words (maximum 32), consistent with Twitter's character limit. Key Twitter-specific features include cashtags (\$AAPL, \$TSLA), hashtags (#earnings, #bearish), mentions (@analyst) and URLs. Bullish tweets contain more cashtags on average (0.34 vs. 0.22 for Bearish), suggesting that positive sentiment correlates with specific stock discussion. The most discriminative terms per class are: **Bearish** — 'down', 'miss', 'cut', 'bearish'; **Bullish** — 'beats', 'strong', 'upgraded', 'buy'; **Neutral** — 'flat', 'unchanged', 'stable'. Per-class word clouds are included in the accompanying notebook.

2.3 Text-Quality Audit: Mojibake and Truncation

A character-level audit reveals that the corpus is *corrupted at source*: 7.0% of training tweets contain UTF-8 mojibake (byte sequences decoded with the wrong codec, e.g. â instead of an apostrophe), 16.4% end in truncation artefacts (. . ./replacement characters cutting the sentence mid-token), and 46.8% carry URLs with no sentiment value. In total, 50.1% of tweets are altered by our repair step. Corruption is not uniform across classes (Neutral 8.7%, Bearish 5.5%, Bullish 2.9%), so leaving it unfixed injects class-correlated noise. This finding motivates the `fix_text` preprocessing technique (Section 3.1) based on the `ftfy` library (Speer, 2019).

2.4 Lexicon Limits: VADER by Class

Scoring every tweet with VADER (Hutto & Gilbert, 2014) exposes why general-purpose lexicons fail on financial text: among *Neutral* tweets, 32.2% have a positive VADER compound and 18.4% a negative one — over half of the majority class looks polarised to a lexicon. Financial wording (“raised”, “beats”, “cut”, “misses”) is factual news vocabulary, not opinion. This quantifies the need for domain-specific models (FinBERT) over lexicon or generic baselines, and is revisited by the agentic workflow in Section 7, which only trusts VADER when its signal is strong.

2.5 EDA-to-Modelling Decisions

The exploratory analysis is not only descriptive; it directly determines the modelling choices.

Finding	Evidence	Modelling consequence
Severe class imbalance	Neutral 64.7%, Bullish 20.1%, Bearish 15.1%	Stratified folds, class-weighted losses and macro-F1 as the primary metric.
Short, information-dense tweets	Mean 12.2 words, maximum 32 words	Transformer max length 96 covers the corpus with headroom and avoids unnecessary padding.
Financial tokens carry sentiment	Cashtags, analyst actions, earnings and target-price terms	Raw text is preserved for the final transformer; aggressive cleaning is tested but not assumed.
URL/mention boilerplate is common	URLs appear in 46.8% of the tweets	Regex cleaning is evaluated for sparse baselines, while tweet-specialised tokenizers handle raw text.
Corpus is corrupted at source	7.0% mojibake, 16.4% truncation artefacts, class-correlated	New <code>fix_text</code> repair step (<code>ftfy</code> + artefact/URL cleanup), ablated for sparse models and transformers.
Lexicons mis-read financial news	50.6% of Neutral tweets look polarised to VADER	Domain-matched encoders (FinBERT) prioritised in the backbone search; VADER demoted to agent baseline.

Table 1: How EDA findings are translated into concrete modelling choices.

2.6 Corpus Split and Validation Protocol

To satisfy the corpus-split requirement without wasting labelled data, every experiment uses **stratified k -fold cross-validation** on the 9,543 labelled tweets (`StratifiedKFold`, `shuffle=True`, `random_state=42`). Stratification keeps the Bearish/Bullish/Neutral proportions stable in each validation fold, which is essential under the 4.28:1 imbalance. All transformer fine-tuning and the final solution use **10 folds** (train on 90%, validate on 10%), which both enlarges each training split and stabilises the averaged test prediction; the broad classical-ML comparison (68 model×feature combinations, Section 5.1) uses 5 folds for tractability. Out-of-fold (OOF) predictions give honest, leak-free macro-F1 on the full corpus, and the fold-level test probability vectors are averaged to produce `pred_33.csv`. The identical seed and splitter make every OOF score in this report directly

comparable within its protocol.

3 Data Preprocessing

We implement six preprocessing techniques following the NLTK-based approach from course Lab 1 and the tweet-tokenisation practice from the extra TF-IDF lab. Crucially, preprocessing is treated as an *empirical question*: every variant is measured under the same 5-fold protocol before deciding what enters the final pipeline.

3.1 Preprocessing Pipeline

- **T1 — Twitter noise cleaning.** Regex removal of URLs and mentions ($\rightarrow$ USER); hashtag-symbol stripping; cashtags mapped to TICKER_X.
- **T2 — Unicode normalisation.** NFKD decomposition then ASCII encoding.
- **T3 — Stop-word removal.** NLTK English list, negations preserved ({not, no, never, neither, nor, none}) — critical for sentiment (Manning et al., 2008).
- **T4 — WordNet lemmatisation.** POS-aware: ‘beaten’ $\rightarrow$ ‘beat’, ‘falling’ $\rightarrow$ ‘fall’.
- **T5 — Snowball stemming.** More aggressive: ‘earnings’ $\rightarrow$ ‘earn’.
- **T6 — TweetTokenizer.** `reduce_len=True`, `strip_handles=True`.
- **T7 — fix_text (encoding repair).** Motivated by the text-quality audit (Section 2.3): `ftfy` mojibake repair (Speer, 2019), removal of end-of-tweet truncation artefacts and of sentiment-free URLs. Unlike T1–T6, this step *restores* signal rather than removing it, and alters 50.1% of tweets.

Example transformations are shown in the notebook: URLs and mentions become generic tokens, cashtags retain the ticker identity, negations survive stop-word removal, and lemmatisation/stemming reduce inflected forms. This demonstrates compliance with the preprocessing requirement while also exposing the main trade-off: cleaning removes noise, but may also remove finance-specific signal.

3.2 Ablation Study

Each configuration is evaluated with Logistic Regression ($C = 1.0$) and TF-IDF (10,000 vocabulary) under 5-fold stratified CV.




Figure 2: Preprocessing ablation: macro-F1 by configuration (green = best). Encoding repair (`fix_text`) edges out raw text; every destructive pipeline hurts.

Preprocessing configuration	Macro-F1	Std
raw + fix_text (ftfy repair)	0.7186	0.0028
raw (no preprocessing)	0.7183	0.0053
clean + stop-words + stemming	0.7118	0.0089
clean + stop-words + lemmatisation	0.7086	0.0042
clean only	0.7024	0.0113

Table 2: Preprocessing ablation (5-fold CV, LR + TF-IDF). Minimal intervention wins: encoding repair (shaded) edges out raw text, while every aggressive pipeline hurts.

Key finding: aggressive preprocessing *destroys* signal — ticker symbols (\$AAPL), analyst-firm names and domain vocabulary ('downgrade', 'EPS beat', 'guidance raise') are discriminative features (Go et al., 2009) — whereas *repairing* the text helps. For sparse TF-IDF the fix_text gain is small (+0.03pp, the vectoriser tokenises most artefacts away), but for fine-tuned transformers, whose subword vocabularies are sensitive to corrupted bytes, the same repair yields **+0.26pp** on FinBERT (0.9056 → 0.9082) and +0.07pp on DeBERTa-v3-large (0.8882 → 0.8889) under the 10-fold protocol (full ablation in `tm_tests_33.ipynb`, Section 5g). The final pipeline therefore feeds the transformer *repaired raw text*: fix_text only, no destructive cleaning.

4 Feature Engineering

Eleven feature representations are computed, covering the three families required by the brief: bag-of-words, Word2Vec and Transformer Encoder.

4.1 Sparse Representations (BoW)

Four sparse sets on raw text: **BoW** (binary CountVectorizer, `max_features=20k`, `min_df=2`); **TF-IDF 1g** (`max_features=20k`, `max_df=0.80`, `min_df=2`); **TF-IDF 2g** (uni+bigram, `max_features=50k`); and **TF-IDF char** (char 3–5-gram, `max_features=30k`).

4.2 Word Embeddings (Word2Vec / GloVe)

Word2Vec Skip-gram and **CBOW** (200d) trained on the corpus (Mikolov et al., 2013), pooled by TF-IDF-weighted mean; and **GloVe-Twitter-100** pre-trained on 2B tweets (Pennington et al., 2014).

4.3 Transformer Embeddings (Encoder)

Three transformer representations (CLS / mean-pooling): **FinBERT** (Araci, 2019); **SBERT all-mpnet-base-v2** (Reimers & Gurevych, 2019; +0.50 extra); and **Twitter-RoBERTa** (Barbieri et al., 2020; +0.50 extra). The latter two count as additional Transformer Encoder methods for the extra-work criterion, and the Twitter-RoBERTa family is carried forward to end-to-end fine-tuning in Section 5.2.

4.4 Financial Hand-crafted Features

A 14-dimensional vector combines VADER scores (Hutto & Gilbert, 2014) with surface statistics (cashtag/hashtag/mention/URL counts, length, capitalisation, punctuation) and a Flesch–Kincaid readability score.

5 Classification Models

All models are evaluated via stratified cross-validation (`StratifiedKFold`, `random_state=42`) with macro-F1 as the primary metric — 5 folds for the broad classical sweep (Section 5.1), 10 folds for all transformer fine-tuning and the final solution (Sections 5.2–5.4). Within each protocol the identical split makes the out-of-fold (OOF) scores directly comparable.

5.1 Traditional ML Models

- **Naïve Bayes.** MultinomialNB ($\alpha = 1.0$) and ComplementNB ($\alpha = 0.5$) on BoW / TF-IDF; ComplementNB targets imbalance (Rennie et al., 2003).
- **Logistic Regression.** L_2 (SAGA), $C \in \{0.1, 1.0, 10.0\}$, all feature sets.
- **Linear SVM.** CalibratedClassifierCV(LinearSVC), $C \in \{0.1, 1.0\}$.
- **KNN.** $k \in \{5, 15\}$, cosine distance, dense embeddings.
- **MLP.** (256, 128) and (512, 256, 128), ReLU, early stopping.
- **Random Forest / XGBoost / LightGBM.** 100 and 300 estimators; LightGBM is the strongest traditional learner on dense embeddings.

5.2 Transformer Encoder Fine-tuning + Backbone Search

Whereas Section 4.3 used the encoders only as *frozen* embeddings, here the full encoder is fine-tuned end-to-end (Devlin et al., 2019) — the single largest lever in this study. Using the Hugging Face Hub for **model selection**, we shortlist nine candidate configurations and fine-tune each under the identical 10-fold protocol: 5–10 epochs, learning rate 5×10^{-6} – 10^{-5} with cosine warmup, batch 8–16 (gradient accumulation to effective 32 for the large models), max length 128, AdamW, gradient clipping 1.0, inverse-frequency class-weighted cross-entropy with label smoothing 0.05, fp16 mixed precision with loss scaling, and best-checkpoint-per-fold selection on validation F1 (NVIDIA RTX 5070). Table 3 reports the comparison.

Backbone (fine-tuned end-to-end, 10-fold)	Params	OOF F1	Acc.
FinBERT-fintwitter, 10 ep + fix_text	110M	0.9082	0.9293
FinBERT-fintwitter, 7 ep + LLRD 0.9 + fix_text	110M	0.9080	0.9285
FinBERT-fintwitter, 10 ep (raw text)	110M	0.9056	0.9275
FinBERT-fintwitter, 7 ep	110M	0.9049	0.9264
FinBERT-fintwitter, 5 ep	110M	0.9025	0.9243
twitter-roberta-large-topic-sentiment (prev. submission)	355M	0.8943	0.9141
DeBERTa-v3-large, 6 ep + fix_text (He et al., 2021) [†]	435M	0.8889	0.9107
DeBERTa-v3-large, 6 ep	435M	0.8882	0.9100
DeBERTa-v3-base finance, 8 ep + fix_text	184M	0.8715	0.8947

Table 3: Transformer backbone search (identical 10-fold protocol, seed 42). The shaded FinBERT-fintwitter is the best individual encoder and the distillation student of Section 5.3. LLRD = layer-wise learning-rate decay. [†]DeBERTa-v3 requires fp32 master weights with loss scaling — its disentangled attention overflows in plain fp16.

The decisive factor is **domain-matched pre-training, not size**: the winning encoder (nickmuchi/finbert-tone-finetuned-fintwitter-classification) is a 110M-parameter FinBERT (Araci, 2019) already adapted to financial-Twitter text, and it outperforms a four-times-larger DeBERTa-v3-large by ~ 2 pp. We replace its classification head with a fresh 3-class head and fine-tune the complete network. Best individual configuration: OOF macro-F1 **0.9082** (± 0.0098),

accuracy 0.9293, precision 0.9008, recall 0.9163 — a **+0.1061** gain over the 0.8021 frozen-embedding baseline and +0.0139 over our previous submission candidate (Twitter–RoBERTa-large, 0.8943). Two further observations: (i) the `fix_text` repair contributes +0.26pp at no architectural cost (Section 3.2); (ii) LLRD ties the best model (0.9080) but trains with different layer dynamics — useful diversity for the ensemble below.

5.3 Ensemble and Knowledge Distillation (Extra Work)

The backbone search leaves us with several strong, partially decorrelated encoders. As extra work we combine them by **weighted soft-voting**: test probabilities are averaged with per-model weights found by *exhaustive* OOF grid search — all 2^8 model subsets crossed with weights $\{0.25, 0.5, 0.75, 1.0\}$, evaluated on leak-free OOF predictions. The optimum keeps eight models (four at weight 1.0, three at 0.75, one at 0.5) and reaches OOF macro-F1 **0.9201** — +1.19pp over the best individual encoder. A paired bootstrap over the 9,543 OOF predictions (1,000 resamples) puts the 95% confidence interval of this gain at $[+0.0066, +0.0172]$, excluding zero: the improvement is real, not resampling noise (Section 6.3).

The delivery rules, however, require the final solution to be “a single pipeline with a single classification model” (Guidelines 5.2). Rather than discard the ensemble or break the rule, we **distil** it (Hinton et al., 2015): the eight-model ensemble becomes the *teacher*, and a single FinBERT-fintwitter student is trained on the combined objective

$$\mathcal{L} = (1 - \alpha) \text{CE}_{\text{weighted}}(y, p_s) + \alpha T^2 \text{KL}\left(p_t^{(T)} \parallel p_s^{(T)}\right), \quad \alpha = 0.5, T = 2,$$

where the teacher targets p_t are the ensemble’s *out-of-fold* probabilities — each training sample’s soft target comes from folds that never saw that sample, so the distillation itself is leakage-free. Training otherwise mirrors the 10-fold recipe of Section 5.2. A grid search over the distillation hyperparameters ($\alpha \in \{0.5, 0.7, 0.9, 1.0\}$, $T \in \{2, 3\}$, 3-fold proxy) confirmed $\alpha=0.5$, $T=2$ as optimal; we also evaluated post-hoc per-class probability scaling and *rejected* it, as its held-out gain was within noise — consistent with the calibration analysis of Section 6.3. The distilled student reaches OOF macro-F1 **0.9134** (accuracy 0.9326) — recovering 44% of the ensemble’s gain in a single 110M-parameter model, itself a statistically significant improvement over the best conventionally-trained encoder (paired bootstrap 95% CI: $[+0.0016, +0.0087]$). This is the **submitted solution** (`pred_33.csv`, produced by the single pipeline in `tm_final_33.ipynb`).

5.4 Decoder Model — Fine-tuned GPT-2 (Extra Work, +1.00 pt)

As decoder-based extra work we use a *decoder* (causal) language model for classification. **GPT-2** (Radford et al., 2019) is fine-tuned end-to-end with a sequence-classification head (`GPT2ForSequenceClassification`, which pools the last non-pad token) under the same 5-fold protocol. It reaches an OOF macro-F1 of **0.7724** (accuracy 0.8125) — a genuine decoder-for-classification result. For completeness we also tried *few-shot prompting* of GPT-2 in the in-context-learning style of GPT-3 (Brown et al., 2020), using verbalizer likelihood scoring with contextual calibration (Zhao et al., 2021); it collapses to the majority class (macro-F1 ≈ 0.25), empirically showing that a 124M decoder is too weak to classify zero/few-shot and that fine-tuning is the correct route.

5.5 Hyperparameter Tuning — Optuna

Optuna (Akiba et al., 2019) with a TPE sampler and median pruner runs 50 trials over LightGBM hyperparameters; the best frozen-embedding configuration reaches macro-F1 0.8021 (RoBERTa features). A second study tunes Logistic-Regression C on TF-IDF bigrams; a soft-voting LR+SVC ensemble provides a strong sparse baseline.

5.6 Stacking Ensemble (Extra Work)

As additional work we build a stacked generalisation (Wolpert, 1992): six base learners (Optuna-tuned LightGBM on RoBERTa, MLP on SBERT, XGBoost on RoBERTa, LightGBM on FinBERT, LR on TF-IDF bigrams, and a fine-tuned Twitter-RoBERTa) emit out-of-fold probabilities concatenated into an 18-dimensional meta-feature matrix on which a Logistic-Regression meta-learner is trained. Because every base learner shares the identical fold split, the meta-features are leakage-free. The stack reaches OOF macro-F1 **0.8483** — strong at the time, but ultimately surpassed both by the individually fine-tuned encoders of Section 5.2 and by the soft-vote ensemble of Section 5.3, which combines stronger, more diverse base models.

6 Results and Evaluation

6.1 Model Comparison

Figure 3 ranks the top configurations and Table 4 lists the top models. The eight-encoder ensemble (extra work) tops the ranking at OOF macro-F1 **0.9201**; the **submitted single model** — the FinBERT student distilled from it — follows at **0.9134**, ahead of every conventionally-trained encoder. The progression of the project is visible in the table: frozen embeddings plateau at ~ 0.80 , generic fine-tuned encoders reach 0.88–0.89, the domain-matched FinBERT reaches 0.91, and ensembling plus distillation extracts the last point.



Figure 3: Model comparison under the common OOF protocol (transformers: 10-fold; classical models: 5-fold). Green bars: ensemble teacher and submitted distilled student.

Model	F1 _{mac}	σ	Acc.	Prec.	Rec.
Ensemble, 8 encoders, optimised weights (<i>extra</i>)	0.9201	—	0.9357	0.9097	0.9315
Distilled FinBERT student submitted	0.9134	0.0100	0.9326	0.9017	0.9263
FinBERT-fintwitter 10ep + fix_text	0.9082	0.0098	0.9293	0.9008	0.9163
FinBERT-fintwitter 7ep + LLRD	0.9080	0.0099	0.9285	0.8967	0.9206
Twitter-RoBERTa-large topic-sent. (prev. subm.)	0.8943	0.0100	0.9141	0.8857	0.9036
DeBERTa-v3-large 6ep + fix_text	0.8889	0.0121	0.9107	0.8835	0.8945
Stacking-6 +fine-tuned (<i>extra</i>)	0.8483	0.0073	0.8771	0.8271	0.8757
LightGBM Optuna + RoBERTa embeddings (best classical)	0.8021	—	—	—	—
LightGBM + SBERT + SMOTE (Chawla et al., 2002)	0.7993	—	—	—	—
GPT-2 decoder fine-tuned (<i>extra</i>)	0.7724	0.0159	0.8125	0.7475	0.8171

Table 4: Top models by OOF macro-F1 (transformers: 10-fold; classical/stacking: 5-fold). The shaded row is the **submitted** single model; rows marked (*extra*) are extra work. SMOTE improves the dense-embedding classical baseline by +0.6pp over class weighting but is irrelevant at transformer level, where the class-weighted loss already handles imbalance.

Key findings: (1) dense transformer embeddings outperform sparse TF-IDF and static embeddings across all model families; (2) end-to-end fine-tuning breaks the ~ 0.80 frozen-embedding plateau; (3) within fine-tuning, *domain match beats parameter count* — a 110M FinBERT beats a 435M DeBERTa-v3-large by ~ 2 pp; (4) text repair (fix_text) is a free +0.26pp; (5) ensemble knowledge survives compression into a single student model.

6.2 Error Analysis

Figure 4 shows the confusion matrix of the submitted model. Per class it reaches F1 of **0.871** (Bearish), **0.916** (Bullish) and **0.953** (Neutral). Errors concentrate exactly where the EDA predicted (Section 2.4): the Neutral/directional boundary. Neutral tweets drift into Bullish/Bearish when they report factual but direction-laden news (“price target raised”), and directional tweets get softened into Neutral when the polarity is implicit. Bearish remains the hardest class — it is both the smallest (15.1%) and the most mojibake-affected after Neutral. On the held-out test set the model predicts 393 Bearish, 495 Bullish and 1,500 Neutral (of 2,388) — close to the training prior. Three qualitative failure modes remain: irony/sarcasm, mixed-signal tweets (‘beats EPS but misses revenue’), and truncated tweets whose decisive token was cut off — the ceiling for text-only models on this corpus.




Figure 4: Confusion matrix of the submitted distilled FinBERT (normalised left, absolute counts right), from 10-fold out-of-fold predictions.

6.3 Statistical Significance and Calibration

With only ten fold-level scores, fold-wise tests have low power: a Wilcoxon signed-rank test between the two best individual encoders does not reach significance ($p = 0.065$). We therefore complement it with *paired bootstrap* over the 9,543 OOF predictions (1,000 resamples), which compares models on identical samples:

- ensemble vs. best individual encoder: $\Delta F1$ 95% CI = $[+0.0066, +0.0172]$ — significant;
- distilled student vs. best plain encoder: $[+0.0016, +0.0087]$ — significant;
- best vs. second-best individual encoder: $[-0.0010, +0.0073]$ — *not* significant, confirming that single-encoder differences are within noise and the ensemble/distillation gains are the real levers.

Finally, a temperature-scaling search on the OOF probabilities finds $T=1.0$ optimal for every fine-tuned model: training with label smoothing (0.05) already yields well-calibrated probabilities (confidence-bucket accuracy matches predicted confidence within $\pm 3\text{pp}$ above 0.8), so no post-hoc calibration is applied to the submission.

7 Extra Challenge 2 — Agentic Workflow (+1.50 pts)

We implement a conversational agent that orchestrates three classifiers with *non-trivial*, adaptive routing. The implementation lives in `src/agent.py` and `notebooks/08_Agentic_Workflow.ipynb`, and runs fully offline (a deterministic orchestrator when no LLM key is present, plus an optional LangChain ReAct back-end; Yao et al., 2022).

7.1 Architecture and Routing

The agent exposes three tools — **VADER** (lexical baseline), **LightGBM+SBERT** (ML model) and **FinBERT** (domain transformer) — and routes adaptively:

- VADER gives a quick lexical reading; if $|\text{compound}| > 0.3$ (strong signal) it is trusted directly.
- Otherwise LightGBM+SBERT is consulted; if it agrees with VADER, that is the verdict.
- On *disagreement*, FinBERT is invoked as a domain-expert tie-breaker.
- The agent explains which model it trusted and why, and `ConversationBufferMemory` enables follow-up questions (“why did you classify that as Bearish?”).

This exceeds simple voting through signal-strength-based tool selection, explicit disagreement detection with principled tie-breaking, and transparent reasoning. We validated the workflow end-to-end: on a held-out sample it produces correct, explained verdicts (e.g. “\$TSLA misses estimates, downgrades” → Bearish via FinBERT tie-break), and it is deterministic and reproducible for the oral defence.

8 Conclusion

This work delivers a complete NLP pipeline for financial tweet sentiment classification. Our **submitted single-model solution** — one FinBERT encoder trained by **knowledge distillation** from an eight-encoder ensemble teacher — achieves an out-of-fold macro-F1 of **0.9134** (accuracy 0.9326), a **+0.111** gain over the 0.8021 classical baseline, while remaining a single 110M-parameter classifier. Key contributions: (1) a text-quality audit showing the corpus is corrupted at source, and a repair step (`ftfy fix_text`) worth +0.26pp at transformer level — whereas *destructive* preprocessing consistently hurts; (2) systematic evaluation of 11 feature representations and 12 model families under a common seeded protocol; (3) a nine-configuration transformer backbone search showing

that *domain-matched pre-training beats parameter count* (110M FinBERT > 435M DeBERTa-v3-large by $\sim 2\text{pp}$); (4, extra) a weighted soft-vote ensemble (0.9201) whose gain is statistically significant, compressed into the submitted student with 44% of the gain retained; (5, extra) a fine-tuned GPT-2 *decoder* classifier (0.7724) with the finding that few-shot prompting of a small decoder collapses to the majority class; (6, extra) an explainable, reproducible agentic workflow. Future work could explore label-noise auditing of the Neutral boundary, domain-adaptive pre-training on a larger financial-tweet corpus, and QLoRA fine-tuning of decoder LLMs.

References

- Akiba, T., Sano, S., Yanase, T., Ohta, T., & Koyama, M. (2019). Optuna: A next-generation hyperparameter optimization framework. *KDD 2019*.
- Alammar, J., & Grootendorst, M. (2024). *Hands-On Large Language Models: Language Understanding and Generation*. O'Reilly Media. [Course textbook.]
- Antweiler, W., & Frank, M. Z. (2004). Is all that talk just noise? The information content of internet stock message boards. *Journal of Finance*, 59(3), 1259–1294.
- Araci, D. (2019). FinBERT: Financial sentiment analysis with pre-trained language models. *arXiv:1908.10063*.
- Barbieri, F., Camacho-Collados, J., Espinosa Anke, L., & Neves, L. (2020). TweetEval: Unified benchmark and comparative evaluation for tweet classification. *EMNLP-Findings 2020*.
- Bollen, J., Mao, H., & Zeng, X. J. (2011). Twitter mood predicts the stock market. *Journal of Computational Science*, 2(1), 1–8.
- Brown, T. B., et al. (2020). Language models are few-shot learners. *NeurIPS 2020*.
- Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: Synthetic minority over-sampling technique. *JAIR*, 16, 321–357.
- Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. *NAACL-HLT 2019*.
- Go, A., Bhayani, R., & Huang, L. (2009). Twitter sentiment classification using distant supervision. *Stanford CS224N Report*.
- Hinton, G., Vinyals, O., & Dean, J. (2015). Distilling the knowledge in a neural network. *arXiv:1503.02531*.
- He, P., Gao, J., & Chen, W. (2021). DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing. *arXiv:2111.09543*.
- Hutto, C. J., & Gilbert, E. (2014). VADER: A parsimonious rule-based model for sentiment analysis of social media text. *ICWSM 2014*.
- Loureiro, D., Barbieri, F., Neves, L., Espinosa Anke, L., & Camacho-Collados, J. (2022). TimeLMs: Diachronic language models from Twitter. *ACL 2022 (System Demonstrations)*; *arXiv:2202.03829*.
- Manning, C. D., Raghavan, P., & Schütze, H. (2008). *Introduction to Information Retrieval*. Cambridge University Press.
- Mikolov, T., Chen, K., Corrado, G., & Dean, J. (2013). Efficient estimation of word representations in vector space. *arXiv:1301.3781*.
- Pennington, J., Socher, R., & Manning, C. D. (2014). GloVe: Global vectors for word representation. *EMNLP 2014*.
- Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., & Sutskever, I. (2019). Language models are unsupervised multitask learners. *OpenAI Blog*.
- Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence embeddings using Siamese BERT-networks. *EMNLP 2019*.
- Rennie, J. D., Shih, L., Teevan, J., & Karger, D. (2003). Tackling the poor assumptions of Naive Bayes text

classifiers. *ICML 2003*.

Sebastiani, F. (2002). Machine learning in automated text categorization. *ACM Computing Surveys*, 34(1), 1–47.

Speer, R. (2019). ftfy: Fixes text for you (Version 5.5). *Zenodo*. doi:10.5281/zenodo.2591652.

Wolpert, D. H. (1992). Stacked generalization. *Neural Networks*, 5(2), 241–259.

Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2022). ReAct: Synergizing reasoning and acting in language models. *arXiv:2210.03629*.

Zhao, T. Z., Wallace, E., Feng, S., Klein, D., & Singh, S. (2021). Calibrate before use: Improving few-shot performance of language models. *ICML 2021*.